

loan-approval-project

December 23, 2023

1 1. Importing libraries and dataset

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score
```

```
[2]: df = pd.read_csv('loan.csv')
```

2 2. Data description

```
[3]: df.columns
```

```
[3]: Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
        'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
        'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],
        dtype='object')
```

```
[4]: df.head()
```

```
[4]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	

4	6000	0.0	141.0	360.0
---	------	-----	-------	-------

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID                614 non-null   object
1   Gender                 601 non-null   object
2   Married                 611 non-null   object
3   Dependents              599 non-null   object
4   Education               614 non-null   object
5   Self_Employed           582 non-null   object
6   ApplicantIncome         614 non-null   int64
7   CoapplicantIncome       614 non-null   float64
8   LoanAmount              592 non-null   float64
9   Loan_Amount_Term        600 non-null   float64
10  Credit_History          564 non-null   float64
11  Property_Area           614 non-null   object
12  Loan_Status             614 non-null   object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB
```

```
[6]: df.describe()
```

```
[6]:
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term \
count	614.000000	614.000000	592.000000	600.000000
mean	5403.459283	1621.245798	146.412162	342.000000
std	6109.041673	2926.248369	85.587325	65.12041
min	150.000000	0.000000	9.000000	12.000000
25%	2877.500000	0.000000	100.000000	360.000000
50%	3812.500000	1188.500000	128.000000	360.000000
75%	5795.000000	2297.250000	168.000000	360.000000
max	81000.000000	41667.000000	700.000000	480.000000

	Credit_History
count	564.000000
mean	0.842199

```
std      0.364878
min      0.000000
25%      1.000000
50%      1.000000
75%      1.000000
max      1.000000
```


```
[7]: df.isnull().sum()
```

```
[7]: Loan_ID      0
     Gender      13
     Married      3
     Dependents   15
     Education     0
     Self_Employed 32
     ApplicantIncome 0
     CoapplicantIncome 0
     LoanAmount    22
     Loan_Amount_Term 14
     Credit_History 50
     Property_Area  0
     Loan_Status    0
     dtype: int64
```


3 Exploratory data analysis

```
[8]: def plot(col):
     sns.countplot(x = col, hue = 'Loan_Status', data = df)
```

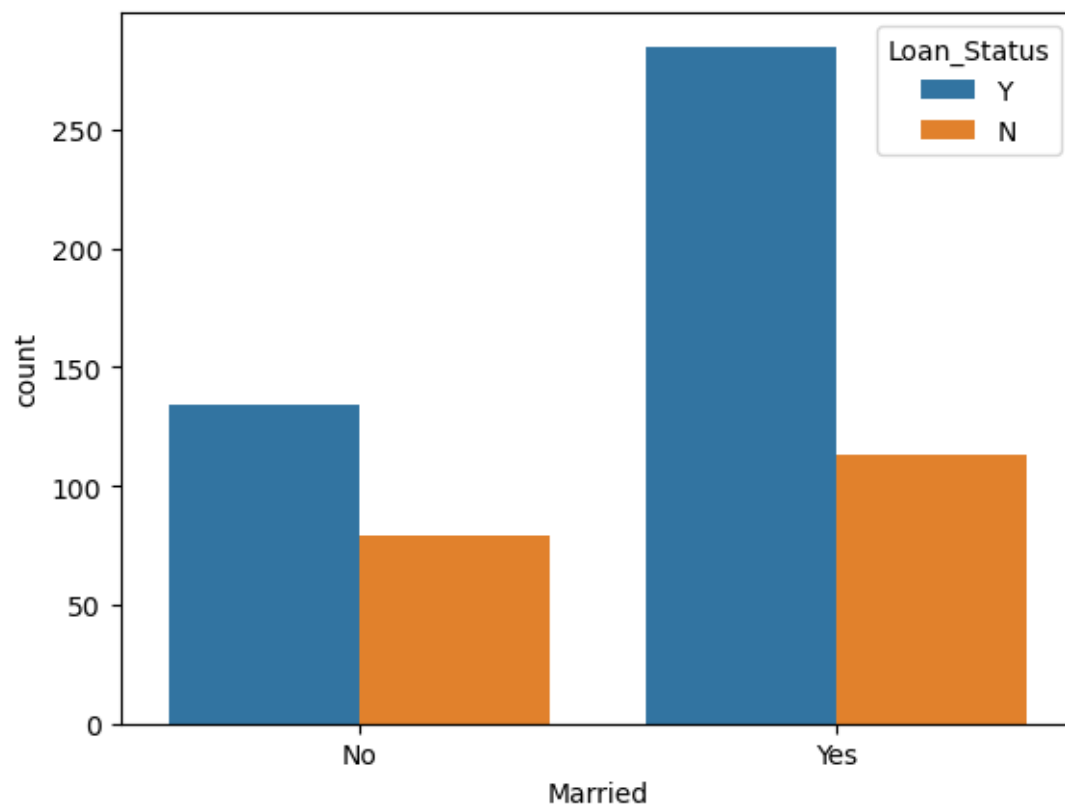
```
[9]: plot('Gender')
```




```
[10]: plot('Education')
```



```
[11]: plot('Married')
```



```
[12]: plot('Credit_History')
```



3.1 4. Dealing with missing values

```
[13]: df.isnull().sum()*100/len(df)
```

```
[13]: Loan_ID          0.000000
      Gender          2.117264
      Married         0.488599
      Dependents      2.442997
      Education        0.000000
      Self_Employed    5.211726
      ApplicantIncome  0.000000
      CoapplicantIncome 0.000000
      LoanAmount       3.583062
      Loan_Amount_Term  2.280130
      Credit_History    8.143322
      Property_Area     0.000000
      Loan_Status       0.000000
      dtype: float64
```

```
[14]: #filling the data for features having missing values > 5%
df['Self_Employed'].fillna(df['Self_Employed'].mode()[0], inplace = True)
df['Credit_History'].fillna(df['Credit_History'].mode()[0], inplace = True)
```

```
[15]: #dropping data for features with missing values < 5%
df.dropna(inplace = True)
```


```
[16]: df.isnull().sum()
```

```
[16]: Loan_ID          0
      Gender         0
      Married        0
      Dependents     0
      Education      0
      Self_Employed  0
      ApplicantIncome 0
      CoapplicantIncome 0
      LoanAmount     0
      Loan_Amount_Term 0
      Credit_History 0
      Property_Area   0
      Loan_Status     0
      dtype: int64
```

3.2 Normalizing the data


```
[17]: sns.kdeplot(df.ApplicantIncome)
```

```
[17]: <Axes: xlabel='ApplicantIncome', ylabel='Density'>
```


```
[18]: sns.kdeplot(df.CoapplicantIncome)
```

```
[18]: <Axes: xlabel='CoapplicantIncome', ylabel='Density'>
```




```
[19]: sns.kdeplot(df.LoanAmount)
```

```
[19]: <Axes: xlabel='LoanAmount', ylabel='Density'>
```




```
[20]: sns.kdeplot(df.Loan_Amount_Term)
```

```
[20]: <Axes: xlabel='Loan_Amount_Term', ylabel='Density'>
```



```
[21]: Total_income = df.ApplicantIncome + df.CoapplicantIncome
      # we made this new feature because loan approval depends on the total income of
      ↳ the family.
      sns.kdeplot(Total_income)
```

```
[21]: <Axes: ylabel='Density'>
```



```
[22]: # as all above drawn graphs are either right skewed or left skewed
      #so we need to normalize the data in order to get more accurate results.
```

```
[23]: Log_loanamount = np.log(df.LoanAmount)
      Log_loanamountterm = np.log(df.Loan_Amount_Term)
      Log_totalincome = np.log(Total_income)
```

```
[24]: df = df.assign(log_totalincome = Log_totalincome)
      df = df.assign(log_loanamount = Log_loanamount)
      df = df.assign(log_loanamountterm = Log_loanamountterm)
```

```
[25]: df.columns
```

```
[25]: Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
          'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
          'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status',
          'log_totalincome', 'log_loanamount', 'log_loanamountterm'],
          dtype='object')
```

```
[26]: col = ['ApplicantIncome', 'CoapplicantIncome', 'Loan_Amount_Term',
          ↪ 'LoanAmount', 'Loan_ID']
      df = df.drop(col, axis = 1)
```

```
[27]: df.columns
```

```
[27]: Index(['Gender', 'Married', 'Dependents', 'Education', 'Self_Employed',  
         'Credit_History', 'Property_Area', 'Loan_Status', 'log_totalincome',  
         'log_loanamount', 'log_loanamountterm'],  
        dtype='object')
```

3.3 Converting data into array

```
[28]: col = ['Loan_Status']  
x = df.drop(col, axis = 1)  
y = df.Loan_Status
```

```
[29]: X = x.iloc[:,:].values  
Y = np.array(y)
```

```
[30]: X
```

```
[30]: array([[ 'Male', 'Yes', '1', ..., 8.714567550836485, 4.852030263919617,  
         5.886104031450156],  
         [ 'Male', 'Yes', '0', ..., 8.006367567650246, 4.189654742026425,  
         5.886104031450156],  
         [ 'Male', 'Yes', '0', ..., 8.50532301884575, 4.787491742782046,  
         5.886104031450156],  
         ...,  
         [ 'Male', 'Yes', '1', ..., 9.025455532779063, 5.53338948872752,  
         5.886104031450156],  
         [ 'Male', 'Yes', '2', ..., 8.933664178700935, 5.231108616854587,  
         5.886104031450156],  
         [ 'Female', 'No', '0', ..., 8.430109084509125, 4.890349128221754,  
         5.886104031450156]], dtype=object)
```

```
[31]: X.shape
```

```
[31]: (553, 10)
```

```
[32]: Y
```

```
[32]: array([ 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'N', 'Y', 'Y', 'Y', 'N',  
         'Y', 'Y', 'Y', 'N', 'N', 'N', 'Y', 'N', 'N', 'Y', 'Y', 'Y', 'N',  
         'Y', 'N', 'N', 'N', 'Y', 'N', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'Y',  
         'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'N', 'N', 'N', 'Y', 'Y', 'N', 'Y',  
         'Y', 'Y', 'Y', 'N', 'N', 'N', 'N', 'Y', 'Y', 'N', 'Y', 'Y', 'Y',  
         'Y', 'N', 'N', 'N', 'N', 'Y', 'Y', 'N', 'N', 'Y', 'Y', 'Y', 'Y',  
         'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',  
         'Y', 'Y', 'N', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'N', 'Y',  
         'Y', 'Y', 'Y', 'Y', 'Y', 'N', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'N',
```

```

'N', 'Y', 'N', 'N', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'N',
'Y', 'N', 'Y', 'N', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'N',
'N', 'Y', 'Y', 'N', 'Y', 'N', 'Y', 'Y', 'N', 'Y', 'N', 'Y', 'Y',
'N', 'Y', 'N', 'N', 'N', 'Y', 'N', 'Y', 'Y', 'N', 'Y', 'Y', 'Y',
'N', 'N', 'Y', 'Y', 'N', 'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'Y',
'Y', 'Y', 'Y', 'N', 'N', 'N', 'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'N',
'Y', 'N', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
'N', 'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'Y',
'Y', 'N', 'N', 'Y', 'Y', 'N', 'Y', 'N', 'N', 'N', 'N', 'Y', 'Y',
'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
'Y', 'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'Y',
'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'N', 'N', 'N', 'Y', 'Y',
'Y', 'Y', 'N', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
'N', 'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'Y',
'Y', 'Y', 'Y', 'N', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
'N', 'Y', 'N', 'Y', 'Y', 'N', 'N', 'Y', 'N', 'Y', 'Y', 'Y', 'Y',
'Y', 'N', 'Y', 'N', 'Y', 'N', 'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'Y',
'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'N', 'N', 'Y', 'N', 'Y', 'Y',
'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'N', 'Y',
'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'N', 'N',
'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'N',
'Y', 'N', 'Y', 'Y', 'N', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
'Y', 'Y', 'N', 'Y', 'Y', 'N', 'N', 'N', 'Y', 'Y', 'N', 'N', 'Y',
'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'N', 'N', 'N', 'Y', 'Y', 'N', 'Y',
'N', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'Y',
'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'N'] , dtype=object)

```

3.4 Label encoding

```
[33]: le = LabelEncoder()
```

```
[34]: for i in range(0,10):
      X[:,i] = le.fit_transform(X[:,i])
```

```
[35]: Y = le.fit_transform(Y)
```

3.5 Feature scaling

```
[36]: ss = StandardScaler()  
      X = ss.fit_transform(X)
```

```
[37]: X
```

```
[37]: array([[ 0.48127522,  0.73511222,  0.23330784, ...,  0.37019928,  
            -0.11941535,  0.30417021],  
          [ 0.48127522,  0.73511222, -0.75914782, ..., -1.48885266,  
            -1.32087679,  0.30417021],  
          [ 0.48127522,  0.73511222, -0.75914782, ..., -0.24714627,  
            -0.29417338,  0.30417021],  
          ...,  
          [ 0.48127522,  0.73511222,  0.23330784, ...,  1.00157541,  
            1.62816493,  0.30417021],  
          [ 0.48127522,  0.73511222,  1.2257635 , ...,  0.87530019,  
            0.97282232,  0.30417021],  
          [-2.0778132 , -1.36033653, -0.75914782, ..., -0.56283433,  
            -0.01019158,  0.30417021]])
```

```
[38]: Y[Y==0]==-1
```

```
[39]: Y
```

```
[39]: array([-1,  1,  1,  1,  1,  1, -1,  1, -1,  1,  1,  1, -1,  1,  1,  1, -1,  
            -1, -1,  1, -1, -1,  1,  1,  1, -1,  1, -1, -1, -1,  1, -1,  1,  1,  
            1, -1,  1,  1,  1,  1,  1, -1,  1,  1,  1, -1, -1, -1,  1,  1, -1,  
            1,  1,  1,  1, -1, -1, -1, -1,  1,  1, -1,  1,  1,  1,  1, -1, -1,  
            -1, -1,  1,  1, -1, -1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  
            1,  1,  1,  1,  1,  1,  1,  1, -1, -1,  1,  1,  1,  1,  1,  1,  1,  
            -1,  1,  1,  1,  1,  1,  1, -1, -1,  1,  1,  1,  1,  1, -1, -1,  1,  
            -1, -1, -1,  1,  1,  1,  1,  1,  1,  1, -1,  1, -1,  1, -1, -1,  1,  
            1,  1,  1,  1,  1,  1, -1, -1,  1,  1, -1,  1, -1,  1,  1, -1,  1,  
            -1,  1,  1, -1,  1, -1, -1, -1,  1, -1,  1,  1, -1,  1,  1,  1, -1,  
            -1,  1,  1, -1,  1,  1, -1,  1,  1,  1,  1,  1,  1,  1,  1, -1, -1,  
            -1,  1,  1,  1,  1, -1,  1, -1,  1, -1,  1,  1,  1, -1,  1,  1,  1,  
            1, -1, -1,  1,  1, -1,  1, -1, -1, -1, -1,  1,  1, -1,  1,  1,  1,  
            1,  1, -1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1, -1,  1,  
            1,  1, -1,  1,  1,  1,  1, -1,  1,  1,  1,  1,  1,  1, -1, -1, -1,  1,  
            1,  1,  1, -1, -1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1, -1,  1,  1,  
            1,  1, -1,  1,  1, -1,  1,  1,  1,  1,  1,  1,  1,  1,  1, -1, -1,  1,  
            1,  1,  1,  1,  1,  1, -1,  1, -1,  1,  1, -1, -1,  1, -1,  1,  1,  
            1,  1, -1, -1,  1, -1,  1,  1,  1, -1,  1,  1,  1,  1,  1, -1,  1,  1,  
            1,  1,  1, -1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1, -1,  1,  1,  
            -1, -1, -1,  1,  1, -1,  1,  1,  1, -1, -1,  1, -1,  1, -1,  1,  1,
```



```
-1, -1, 1, 1, 1, 1, 1, 1, -1, 1, 1, 1, 1, -1, 1, 1, 1,
1, -1, 1, 1, 1, 1, 1, 1, 1, 1, -1, -1, -1, -1, 1, -1, 1,
1, 1, 1, -1, 1, -1, 1, 1, 1, -1, 1, -1, 1, -1, 1, -1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, -1, -1, 1, -1, 1, 1,
1, 1, -1, 1, 1, 1, 1, -1, 1, 1, 1, -1, 1, 1, 1, 1, 1,
-1, 1, 1, -1, -1, 1, 1, -1, -1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, -1, -1, 1, 1, 1, 1, -1, 1, 1, -1, 1, 1, 1, 1,
-1, 1, -1, 1, 1, -1, -1, 1, 1, 1, 1, 1, 1, 1, 1, 1, -1,
1, 1, -1, -1, -1, 1, 1, -1, -1, 1, 1, 1, 1, 1, -1, 1, -1,
-1, -1, 1, 1, -1, 1, -1, 1, 1, 1, -1, 1, 1, -1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, -1])
```

3.6 Splitting data into train and test data

```
[40]: X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size = 0.25,
↳random_state = 7)
```

```
[41]: X_train
```

```
[41]: array([[ 0.48127522, -1.36033653, -0.75914782, ..., -1.6782655 ,
-1.3645663 ,  1.18483055],
[ 0.48127522, -1.36033653, -0.75914782, ...,  1.31024819,
 1.23495936,  0.30417021],
[ 0.48127522,  0.73511222, -0.75914782, ..., -0.24714627,
-0.29417338,  0.30417021],
...,
[ 0.48127522,  0.73511222,  1.2257635 , ...,  0.80514728,
 0.97282232,  0.30417021],
[ 0.48127522,  0.73511222, -0.75914782, ..., -1.68528079,
-1.43010056,  0.30417021],
[-2.0778132 ,  0.73511222, -0.75914782, ...,  0.44035218,
-0.44708665,  0.30417021]])
```

```
[42]: X_test
```

```
[42]: array([[ 0.48127522,  0.73511222,  1.2257635 , ..., -0.08579459,
-0.16310486,  0.30417021],
[ 0.48127522,  0.73511222, -0.75914782, ...,  0.2860158 ,
-0.11941535,  0.30417021],
[ 0.48127522, -1.36033653, -0.75914782, ...,  0.55961212,
 0.14272169,  0.30417021],
...,
[-2.0778132 ,  0.73511222, -0.75914782, ..., -0.56283433,
-0.46893141,  0.30417021],
[ 0.48127522,  0.73511222, -0.75914782, ..., -0.75224717,
-0.57815518,  0.30417021],
[ 0.48127522,  0.73511222,  1.2257635 , ...,  1.67504328,
```

1.47525165, 0.30417021]])

[43]: Y_train

```
[43]: array([ -1, -1,  1, -1, -1, -1,  1, -1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,
        -1,  1,  1, -1,  1,  1,  1,  1,  1, -1,  1, -1,  1,  1,  1,  1,  1,  1,
         1,  1, -1, -1,  1,  1, -1,  1, -1, -1,  1,  1,  1, -1, -1,  1,  1,  1,
         1, -1,  1,  1,  1,  1, -1,  1, -1,  1, -1, -1,  1,  1, -1, -1,  1,
         1,  1, -1,  1,  1, -1,  1,  1, -1,  1,  1,  1, -1, -1,  1,  1,  1,
        -1,  1,  1,  1,  1,  1, -1,  1,  1,  1,  1,  1,  1, -1,  1,  1,  1,
         1, -1,  1,  1,  1,  1, -1, -1,  1,  1,  1,  1, -1,  1,  1, -1,  1,
         1,  1,  1,  1,  1,  1, -1,  1,  1,  1, -1,  1,  1,  1, -1,  1,  1,
         1,  1, -1, -1,  1, -1,  1, -1, -1,  1,  1,  1,  1, -1,  1,  1,  1,
         1, -1,  1,  1,  1,  1,  1, -1,  1,  1,  1,  1, -1, -1,  1,  1,  1,
         1,  1,  1,  1,  1, -1,  1,  1,  1, -1, -1,  1,  1, -1, -1,  1,  1,
        -1,  1,  1,  1,  1,  1, -1,  1, -1, -1, -1,  1,  1, -1,  1,  1,  1,
         1,  1,  1,  1, -1, -1, -1,  1, -1,  1,  1,  1, -1,  1, -1,  1, -1,
         1,  1, -1,  1,  1, -1,  1, -1,  1,  1,  1,  1, -1, -1,  1,  1, -1,
         1,  1,  1,  1,  1, -1,  1, -1,  1, -1,  1,  1,  1, -1, -1,  1,  1,
         1,  1,  1, -1,  1,  1,  1, -1,  1,  1,  1,  1,  1,  1, -1,  1,  1,
        -1,  1,  1,  1, -1, -1,  1, -1,  1,  1,  1, -1,  1, -1,  1,  1, -1,
        -1,  1, -1, -1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1, -1,  1,  1,
         1,  1, -1, -1,  1, -1,  1,  1,  1,  1,  1,  1,  1,  1, -1,  1,  1,
        -1,  1, -1, -1,  1, -1,  1,  1,  1, -1, -1,  1, -1,  1,  1,  1, -1,
        -1, -1,  1,  1,  1,  1,  1,  1,  1,  1, -1,  1, -1,  1,  1, -1,  1,
        -1,  1,  1,  1,  1, -1,  1, -1,  1,  1,  1,  1,  1, -1,  1, -1,  1,
         1,  1, -1,  1, -1, -1,  1,  1,  1,  1, -1,  1, -1,  1,  1,  1,  1,
         1, -1,  1, -1, -1,  1,  1, -1,  1,  1, -1,  1, -1,  1,  1,  1, -1,
         1,  1,  1, -1, -1,  1])
```

[44]: Y_test

```
[44]: array([  1,  1,  1,  1,  1, -1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1, -1,
         1, -1, -1,  1,  1,  1,  1,  1, -1,  1, -1, -1,  1,  1,  1,  1, -1,
         1,  1, -1,  1,  1,  1, -1,  1,  1,  1,  1,  1,  1,  1,  1, -1,  1,
         1,  1,  1, -1,  1,  1,  1,  1,  1,  1,  1,  1,  1, -1, -1,  1,  1,
         1,  1,  1,  1,  1,  1,  1,  1,  1,  1, -1,  1, -1,  1,  1, -1,
         1, -1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1, -1,  1,  1,  1,  1,
        -1,  1,  1, -1, -1,  1, -1,  1, -1, -1,  1, -1,  1, -1,  1,  1,  1,
         1, -1,  1, -1, -1,  1,  1,  1,  1,  1,  1, -1,  1,  1, -1,  1, -1,
        -1,  1, -1])
```

3.7 Model training

```
[45]: def classify(model_name,X, Y, classifier):  
      classifier.fit(X_train, Y_train)  
      X_pred = classifier.predict(X_test)  
      Acc = accuracy_score(Y_test, X_pred)  
      print('Accuracy of', model_name, ' is:', Acc*100)  
      score = cross_val_score(classifier, X, Y, cv = 5)  
      print('Cross validation score of', model_name, ' is:', np.mean(score)*100)
```

```
[46]: from sklearn.ensemble import RandomForestClassifier  
      model_name = 'Random Forest Classifier'  
      classifier = RandomForestClassifier()  
      classify(model_name, X, Y, classifier)
```

Accuracy of Random Forest Classifier is: 84.89208633093526
Cross validation score of Random Forest Classifier is: 77.76412776412776

```
[47]: from sklearn.tree import DecisionTreeClassifier  
      model_name = 'Decision Tree Classifier'  
      classifier = DecisionTreeClassifier()  
      classify(model_name, X, Y, classifier)
```

Accuracy of Decision Tree Classifier is: 70.50359712230215
Cross validation score of Decision Tree Classifier is: 69.25798525798525

```
[48]: from sklearn.linear_model import LogisticRegression  
      model_name = 'Logistic Regression'  
      classifier = LogisticRegression()  
      classify(model_name, X, Y, classifier)
```

Accuracy of Logistic Regression is: 84.17266187050359
Cross validation score of Logistic Regression is: 80.47993447993449

```
[49]: from sklearn.naive_bayes import GaussianNB  
      model_name = 'Gaussian Naive Bayes'  
      classifier = GaussianNB()  
      classify(model_name, X, Y, classifier)
```

Accuracy of Gaussian Naive Bayes is: 84.17266187050359
Cross validation score of Gaussian Naive Bayes is: 80.2997542997543

```
[50]: from sklearn.neighbors import KNeighborsClassifier  
      classifier = KNeighborsClassifier()  
      model_name = 'K Neighbors Classifier'  
      classify(model_name, X, Y, classifier)
```

Accuracy of K Neighbors Classifier is: 80.57553956834532

Cross validation score of K Neighbors Classifier is: 77.94430794430795

Based on accuracy and cross val score, we observe that Logistic Regression performs better among all.